

# Python + 머신러닝 예측

- jupyter/pyspark:python3.8.8에서 작성
- 데이터 출처 (캐글 링크 | <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>)

## EDA

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, RobustScaler
from sklearn.manifold import TSNE
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import recall_score, precision_score, f1_score, accuracy_score
```

```
In [2]: df = pd.read_csv('creditcard.csv')
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time    284807 non-null  float64
1    V1       284807 non-null  float64
2    V2       284807 non-null  float64
3    V3       284807 non-null  float64
4    V4       284807 non-null  float64
5    V5       284807 non-null  float64
6    V6       284807 non-null  float64
7    V7       284807 non-null  float64
8    V8       284807 non-null  float64
9    V9       284807 non-null  float64
10   V10      284807 non-null  float64
11   V11      284807 non-null  float64
12   V12      284807 non-null  float64
13   V13      284807 non-null  float64
14   V14      284807 non-null  float64
15   V15      284807 non-null  float64
16   V16      284807 non-null  float64
17   V17      284807 non-null  float64
18   V18      284807 non-null  float64
19   V19      284807 non-null  float64
20   V20      284807 non-null  float64
21   V21      284807 non-null  float64
22   V22      284807 non-null  float64
23   V23      284807 non-null  float64
24   V24      284807 non-null  float64
25   V25      284807 non-null  float64
26   V26      284807 non-null  float64
27   V27      284807 non-null  float64
28   V28      284807 non-null  float64
29   Amount   284807 non-null  float64
30   Class    284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

In [4]:

df.describe()

Out[4]:

	Time	V1	V2	V3	V4	V5	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7

8 rows × 31 columns

```
In [5]: df.head()
```

```
Out[5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

```
In [6]: # Label 값 확인
df['Class'].value_counts()
```

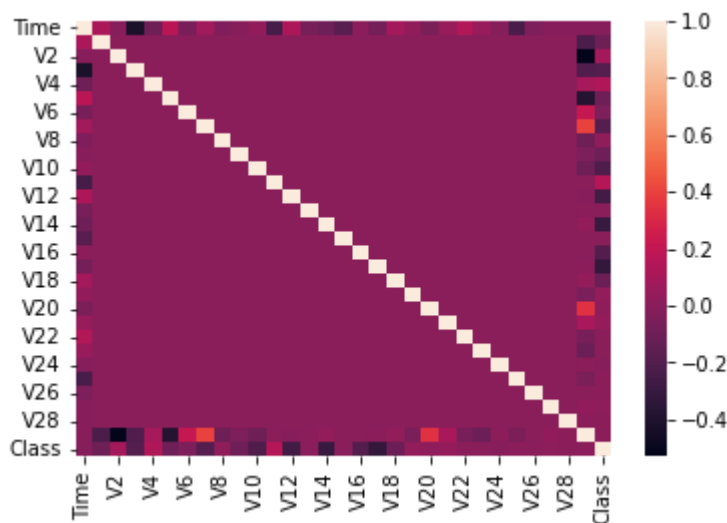
```
Out[6]:
```

0	284315
1	492

Name: Class, dtype: int64

```
In [7]: # 공산성이 낮은 데이터
sns.heatmap(df.corr())
```

```
Out[7]: <AxesSubplot:>
```




```
In [8]: # Label값에 대한 상관관계 확인
df.corr()['Class'].sort_values()
```

```
Out[8]: V17      -0.326481
        V14      -0.302544
        V12      -0.260593
        V10      -0.216883
        V16      -0.196539
        V3       -0.192961
        V7       -0.187257
        V18      -0.111485
        V1       -0.101347
        V9       -0.097733
        V5       -0.094974
        V6       -0.043643
        Time     -0.012323
        V24      -0.007221
        V13      -0.004570
        V15      -0.004223
        V23      -0.002685
        V22       0.000805
        V25       0.003308
        V26       0.004455
        Amount    0.005632
        V28       0.009536
        V27       0.017580
        V8        0.019875
        V20       0.020090
        V19       0.034783
        V21       0.040413
        V2        0.091289
        V4        0.133447
        V11       0.154876
        Class     1.000000
        Name: Class, dtype: float64
```


```
In [9]: # Amount 정규화 필요
sns.displot(data=df, x="Amount", kde=True)
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x7f282cdb6210>
```




```
In [10]: # Time 정규화 필요
sns.displot(data=df, x="Time", kde=True)
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x7f2830d3f0d0>
```



```
In [11]: # standard scaler
std_scaler = StandardScaler()
df['scale_amount_std'] = std_scaler.fit_transform(df['Amount'].values.reshape(-1, 1))
sns.displot(data=df, x="scale_amount_std", kde=True)
```


```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x7f282cd6f450>
```



```
In [12]: # robust scaler
rob_scaler = RobustScaler()
```


```
df['scale_amount_rob'] = rob_scaler.fit_transform(df['Amount'].values.reshape(-1, 1))
sns.displot(data=df, x="scale_amount_rob", kde=True)
```

Out[12]: <seaborn.axisgrid.FacetGrid at 0x7f281d24c550>



In [13]: # Time 정규화  
df['scale\_time\_std'] = std\_scaler.fit\_transform(df['Time'].values.reshape(-1, 1))  
sns.displot(data=df, x="scale\_time\_std", kde=True)

Out[13]: <seaborn.axisgrid.FacetGrid at 0x7f282cc086d0>



In [14]: # 안 쓰는 column 없애기  
df = df.drop(['Time', 'Amount', 'scale\_amount\_std'], axis=1)  
scale\_amount = df['scale\_amount\_rob']  
scale\_time = df['scale\_time\_std']

```
In [15]: df.insert(0, 'scale_amount', scale_amount)
df.insert(1, 'scale_time', scale_time)
df.head()
```

```
Out[15]:
```

	scale_amount	scale_time	V1	V2	V3	V4	V5	V6	V7
0	1.783274	-1.996583	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239595
1	-0.269825	-1.996583	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	4.983721	-1.996562	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791467
3	1.418291	-1.996562	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	0.670579	-1.996541	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592947

5 rows × 33 columns

```
In [16]: # dataset shuffle
df = df.sample(frac=1)
df.head()
```

```
Out[16]:
```

	scale_amount	scale_time	V1	V2	V3	V4	V5	V6	V7
198746	3.884580	0.796372	-0.887797	-0.920315	1.505975	0.054601	-1.032067	0.129266	0.239595
78813	7.621603	-0.781162	0.256881	-2.213319	-0.404803	0.029907	-1.370190	-0.569455	0.791467
236461	-0.194508	1.137299	0.086933	0.995813	-0.329656	-0.626455	0.912197	-0.574713	0.237609
47431	1.229651	-1.086670	-0.931375	0.322760	-0.566347	-0.480548	1.687959	3.933089	-0.078803
283565	-0.148536	1.618978	-1.315703	1.308200	1.583054	0.843591	-0.016332	0.665115	0.592947

5 rows × 33 columns

```
In [17]: fraud_df = df[df['Class'] == 1]
non_fraud_df = df[df['Class'] == 0][:len(fraud_df)]
```

```
In [18]: # undersampling
under_sampled = pd.concat([fraud_df, non_fraud_df])
new_df = under_sampled.sample(frac=1)
new_df.head()
```

Out[18]:

	scale_amount	scale_time	V1	V2	V3	V4	V5	V6	
40336	-0.146720	-1.150118	1.213232	0.019081	0.653910	0.890635	-0.624259	-0.566418	-C
233035	-0.056033	1.107692	1.988135	0.116951	-1.605889	0.351343	0.434941	-0.639378	C
52466	-0.293440	-1.039227	-1.476893	2.122314	-1.229470	1.201849	-0.343264	-1.317704	-1
239214	-0.044435	1.162274	0.143476	0.583583	-0.083261	-0.469798	0.445817	-1.184308	C
95534	0.138476	-0.620279	1.193916	-0.571085	0.742522	-0.014588	-0.624561	0.832162	-C

5 rows × 33 columns

In [19]:

```
print("오리지널 라벨 분포:")
print(df['Class'].value_counts() / len(df))

print("언더 샘플 후 라벨 분포:")
print(new_df['Class'].value_counts() / len(new_df))

sns.countplot('Class', data=new_df)
```

오리지널 라벨 분포:

0 0.998273

1 0.001727

Name: Class, dtype: float64

언더 샘플 후 라벨 분포:

0 0.5

1 0.5


Name: Class, dtype: float64

/opt/conda/envs/py37/lib/python3.7/site-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[19]:

&lt;AxesSubplot:xlabel='Class', ylabel='count'&gt;




In [20]:

```
# 상관관계 확인, 이전과 다르게 상관성을 띠는 피쳐들이 생기기 시작.
under_sample_corr = new_df.corr()
f, ax = plt.subplots(1, figsize=(24,20))
sns.heatmap(under_sample_corr, ax=ax)
```




Out[20]: &lt;AxesSubplot:&gt;



```
In [21]: # 비교
before_corr = df.corr()
f, ax = plt.subplots(1, figsize=(24,20))
sns.heatmap(before_corr, ax=ax)
```

Out[21]: &lt;AxesSubplot:&gt;

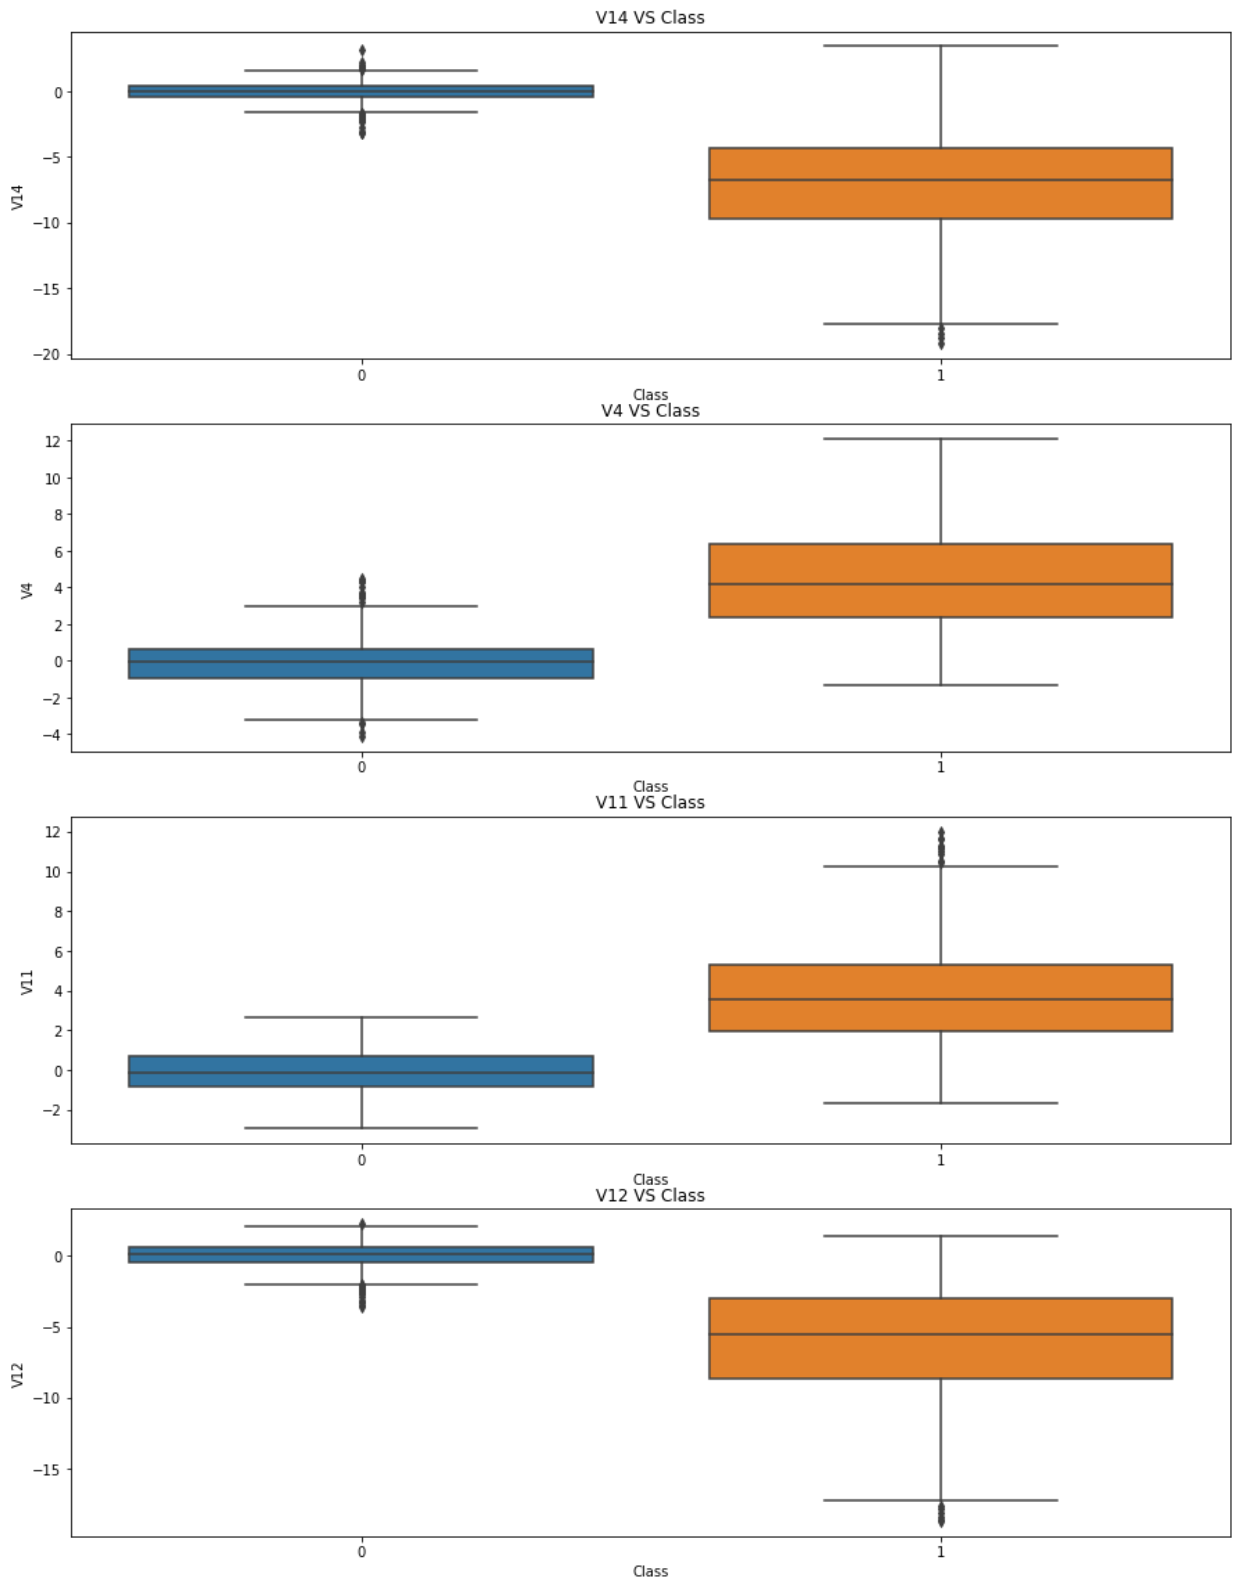


```
In [22]: # class와 가장 상관 관계가 높은 4개 feature 확인
top_4_corr = abs(under_sample_corr['Class']).sort_values(ascending=False)[1:5]
```

```
In [23]: f, ax = plt.subplots(4, figsize=(15, 20))

for i, col in enumerate(top_4_corr.index):
    sns.boxplot(x="Class", y=col, data=new_df, ax=ax[i])
    ax[i].set_title(f"{col} VS Class")

plt.show()
```



```
In [24]: # outlier들이 존재. outlier들을 제거
v14_fraud = new_df['V14'].loc[new_df['Class'] == 1].values
q25 = np.percentile(v14_fraud, 25)
q75 = np.percentile(v14_fraud, 75)
v14_iqr = q75 - q25
v14_cutoff = v14_iqr * 1.5
v14_lower, v14_upper = q25 - v14_cutoff, q75 + v14_cutoff
print('q25: {}, q75: {}'.format(q25, q75))
print('v14_iqr: {}'.format(v14_iqr))
```

```
print('v14_cutoff: {}'.format(v14_cutoff))
print('v14_lower: {} v14_upper: {}'.format(v14_lower, v14_upper))
```

```
q25: -9.692722964972386, q75: -4.282820849486865
v14_iqr: 5.409902115485521
v14_cutoff: 8.114853173228282
v14_lower: -17.807576138200666 v14_upper: 3.8320323237414167
```

```
In [25]: # outliers
outliers = [i for i in v14_fraud if i < v14_lower or i > v14_upper]
outliers
```

```
Out[25]: [-18.8220867423816, -18.0499976898594, -19.2143254902614, -18.4937733551053]
```

```
In [26]: new_df[(new_df['V14'] > v14_upper) | (new_df['V14'] < v14_lower)].index
```

```
Out[26]: Int64Index([8615, 9252, 8296, 9035], dtype='int64')
```

```
In [27]: # v14 outlier 제거
new_df = new_df.drop(new_df[(new_df['V14'] > v14_upper) | (new_df['V14'] < v14_lower)].
```

```
In [28]: # v11 outlier 제거
v11_fraud = new_df['V11'].loc[new_df['Class'] == 1].values
q25 = np.percentile(v11_fraud, 25)
q75 = np.percentile(v11_fraud, 75)

v11_iqr = q75 - q25

v11_cutoff = v11_iqr * 1.5

v11_lower, v11_upper = q25 - v11_cutoff, q75 + v11_cutoff

new_df = new_df.drop(new_df[(new_df['V11'] > v11_upper) | (new_df['V11'] < v11_lower)].
```

```
In [29]: # v12 outlier 제거
v12_fraud = new_df['V12'].loc[new_df['Class'] == 1].values
q25 = np.percentile(v12_fraud, 25)
q75 = np.percentile(v12_fraud, 75)

v12_iqr = q75 - q25

v12_cutoff = v12_iqr * 1.5


v12_lower, v12_upper = q25 - v12_cutoff, q75 + v12_cutoff

new_df = new_df.drop(new_df[(new_df['V12'] > v12_upper) | (new_df['V12'] < v12_lower)].
```

```
In [30]: # outlier 제거 후 시각화
f, ax = plt.subplots(4, figsize=(15, 20))

for i, col in enumerate(top_4_corr.index):
    sns.boxplot(x="Class", y=col, data=new_df, ax=ax[i])
    ax[i].set_title(f"{col} VS Class")

plt.show()
```



```
In [31]: X = new_df.drop('Class', axis=1)
         y = new_df['Class']
```

```
In [32]: x_reduced_tsne = TSNE(n_components=2).fit_transform(X.values)
```

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/manifold/_t_sne.py:783: FutureWarning: The default initialization in TSNE will change from 'random' to 'pca' in 1.2.
FutureWarning,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/manifold/_t_sne.py:793: FutureWarning: The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.
FutureWarning,
```

```
In [33]: tsne_df = new_df.copy()
tsne_df['tsne_1'] = x_reduced_tsne[:,0]
tsne_df['tsne_2'] = x_reduced_tsne[:,1]
```

```
In [34]: # TSNE value 확인
tsne_df.corr()['tsne_2'].sort_values()[:10]
```


```
Out[34]: V14      -0.545031
tsne_1    -0.436242
V12      -0.363767
V6        -0.275328
V16      -0.212727
V10      -0.209848
V9        -0.189780
V24      -0.163188
V17      -0.147153
V3        -0.124907
Name: tsne_2, dtype: float64
```

```
In [35]: # TSNE 시각화
f, ax = plt.subplots(1, figsize=(24,20))

ax.scatter(x_reduced_tsne[:,0], x_reduced_tsne[:,1], c=(y==0), label='Not Fraud')
ax.scatter(x_reduced_tsne[:,0], x_reduced_tsne[:,1], c=(y==1), label='Fraud')

ax.legend()

plt.show()
```



## Model Baseline

```
In [43]: X = df.drop('Class', axis=1)
y = df['Class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [ ]: classifiers = {
    'LR': LogisticRegression(),
    'KM': KNeighborsClassifier(),
    'SVC': SVC(),
    'Decision Tree': DecisionTreeClassifier(),
    'RandomForestClassifier': RandomForestClassifier()
}

X_train = X_train.values
X_test = X_test.values
y_train = y_train.values
y_test = y_test.values

# 모델 정확도 예측
for key, classifier in classifiers.items():
    classifier.fit(X_train, y_train)
    y_pred = cross_val_predict(classifier, X_train, y_train, cv=5)
    precision, recall, threshold = precision_recall_curve(y_train, y_pred)
    print(key)
    print('---' * 20)
```

```
print('Recall Score: {:.2f}'.format(recall_score(y_train, y_pred)))
print('Precision Score: {:.2f}'.format(precision_score(y_train, y_pred)))
print('F1 Score: {:.2f}'.format(f1_score(y_train, y_pred)))
print('Accuracy Score: {:.2f}'.format(accuracy_score(y_train, y_pred)))
print('---' * 20)
#     score = cross_val_score(classifier, X_train, y_train, cv=5)
#     print(key, score.mean())
```



```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

LR

```
-----
Recall Score: 0.62
Precision Score: 0.86
F1 Score: 0.72
Accuracy Score: 1.00
-----
```

## Model

```
In [36]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [37]: X_train = X_train.values
X_test = X_test.values
y_train = y_train.values
y_test = y_test.values
```

```
In [38]: classifiers = {
    'LR': LogisticRegression(),
    'KM': KNeighborsClassifier(),
    'SVC': SVC(),
    'Decision Tree': DecisionTreeClassifier(),
    'RandomForestClassifier': RandomForestClassifier()
}
```

```
In [39]: # 모델 정확도 예측
for key, classifier in classifiers.items():
    classifier.fit(X_train, y_train)
    score = cross_val_score(classifier, X_train, y_train, cv=5)
    print(key, score.mean())
```

```
LR 0.9404692082111439
KM 0.941767909509845
SVC 0.9404692082111437
Decision Tree 0.9223795559279431
RandomForestClassifier 0.9495266024298281
```

```
In [40]: # 모델 튜닝
# Logistic Regression
lr_param = {"penalty": ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}

lr = GridSearchCV(LogisticRegression(), lr_param)
lr.fit(X_train, y_train)
lr_bestparam = lr.best_estimator_

# KNeighbors Classifier
kn_param = {"n_neighbors": list(range(2,5,1)), 'algorithm': ['auto', 'ball_tree', 'kd_

knn = GridSearchCV(KNeighborsClassifier(), kn_param)
knn.fit(X_train, y_train)
knn_bestparam = knn.best_estimator_

# Support Vector Classifier
svc_param = {'C': [0.5, 0.7, 0.9, 1], 'kernel': ['rbf', 'poly', 'sigmoid', 'linear']}
svc = GridSearchCV(SVC(), svc_param)
svc.fit(X_train, y_train)
svc_bestparam = svc.best_estimator_
```

```
# DecisionTree Classifier
tree_param = {"criterion": ["gini", "entropy"], "max_depth": list(range(2,4,1)),
              "min_samples_leaf": list(range(5,7,1))}
dt = GridSearchCV(DecisionTreeClassifier(), tree_param)
dt.fit(X_train, y_train)
dt_bestparam = dt.best_estimator_
```

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,  
 /opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/model\_selection/\_validation.py:372: FitFailedWarning:  
 35 fits failed out of a total of 70.  
 The score on these train-test partitions for these parameters will be set to nan.  
 If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

-----

35 fits failed with the following error:  
 Traceback (most recent call last):  
 File "/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/model\_selection/\_validation.py", line 680, in \_fit\_and\_score  
 estimator.fit(X\_train, y\_train, \*\*fit\_params)  
 File "/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear\_model/\_logistic.py", line 1461, in fit  
 solver = \_check\_solver(self.solver, self.penalty, self.dual)  
 File "/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear\_model/\_logistic.py", line 449, in \_check\_solver  
 % (solver, penalty)  
 ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn(some\_fits\_failed\_message, FitFailedWarning)  
 /opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/model\_selection/\_search.py:972: UserWarning: One or more of the test scores are non-finite: [  
 6 nan 0.94046921 nan 0.94176791 nan 0.9146208  
 nan 0.94046921 nan 0.94306661 nan 0.94436531  
 nan 0.94306661]  
 category=UserWarning,

```
In [41]: # Model accuracy 출력

lr_score = cross_val_score(lr_bestparam, X_train, y_train, cv=5)
print('Logistic Regression Cross Validation Score: ', round(lr_score.mean() * 100, 2)).

knn_score = cross_val_score(knn_bestparam, X_train, y_train, cv=5)
print('Knears Neighbors Cross Validation Score', round(knn_score.mean() * 100, 2)).ast

svc_score = cross_val_score(svc_bestparam, X_train, y_train, cv=5)
print('Support Vector Classifier Cross Validation Score', round(svc_score.mean() * 100, 2)).

dt_score = cross_val_score(dt_bestparam, X_train, y_train, cv=5)
print('DecisionTree Classifier Cross Validation Score', round(dt_score.mean() * 100, 2)).

Logistic Regression Cross Validation Score: 94.44%
Knears Neighbors Cross Validation Score 94.44%
Support Vector Classifier Cross Validation Score 94.83%
DecisionTree Classifier Cross Validation Score 92.5%
```

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
In [42]: lr_pred = cross_val_predict(lr_bestparam, X_train, y_train, cv=5)

precision, recall, threshold = precision_recall_curve(y_train, lr_pred)
y_pred = lr_pred

print('Logistic Regression')
print('---' * 20)
print('Recall Score: {:.2f}'.format(recall_score(y_train, y_pred)))
print('Precision Score: {:.2f}'.format(precision_score(y_train, y_pred)))
print('F1 Score: {:.2f}'.format(f1_score(y_train, y_pred)))
print('Accuracy Score: {:.2f}'.format(accuracy_score(y_train, y_pred)))
print('---' * 20)

knn_pred = cross_val_predict(knn_bestparam, X_train, y_train, cv=5)

precision, recall, threshold = precision_recall_curve(y_train, knn_pred)
y_pred = knn_pred

print('KNeighbors Classifier')
print('---' * 20)
print('Recall Score: {:.2f}'.format(recall_score(y_train, y_pred)))
print('Precision Score: {:.2f}'.format(precision_score(y_train, y_pred)))
print('F1 Score: {:.2f}'.format(f1_score(y_train, y_pred)))
print('Accuracy Score: {:.2f}'.format(accuracy_score(y_train, y_pred)))
print('---' * 20)

svc_pred = cross_val_predict(svc_bestparam, X_train, y_train, cv=5)

precision, recall, threshold = precision_recall_curve(y_train, svc_pred)
y_pred = svc_pred
```

```

print('Support Vector Machine')
print('---' * 20)
print('Recall Score: {:.2f}'.format(recall_score(y_train, y_pred)))
print('Precision Score: {:.2f}'.format(precision_score(y_train, y_pred)))
print('F1 Score: {:.2f}'.format(f1_score(y_train, y_pred)))
print('Accuracy Score: {:.2f}'.format(accuracy_score(y_train, y_pred)))
print('---' * 20)

dt_pred = cross_val_predict(dt_bestparam, X_train, y_train, cv=5)

precision, recall, threshold = precision_recall_curve(y_train, dt_pred)
y_pred = dt_pred

print('Decision Tree')
print('---' * 20)
print('Recall Score: {:.2f}'.format(recall_score(y_train, y_pred)))
print('Precision Score: {:.2f}'.format(precision_score(y_train, y_pred)))
print('F1 Score: {:.2f}'.format(f1_score(y_train, y_pred)))
print('Accuracy Score: {:.2f}'.format(accuracy_score(y_train, y_pred)))
print('---' * 20)

```

Logistic Regression

```

-----
Recall Score: 0.92
Precision Score: 0.96
F1 Score: 0.94
Accuracy Score: 0.94
-----

```

KNeighbors Classifier

```

-----
Recall Score: 0.91
Precision Score: 0.97
F1 Score: 0.94
Accuracy Score: 0.94
-----

```

Support Vector Machine

```

-----
Recall Score: 0.92
Precision Score: 0.97
F1 Score: 0.94
Accuracy Score: 0.95
-----

```

Decision Tree

```

-----
Recall Score: 0.89
Precision Score: 0.95
F1 Score: 0.92
Accuracy Score: 0.92
-----

```

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
8: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
/opt/conda/envs/py37/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:81
```

```
8: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

In [ ]: